



DRL final project

第一組

楊晶宇 伍麒 常安彥

Outline



01

Episodic backward update

A brief introduce of EBU

02

Implementation

Code and hyperparameters

03

Results and Behavior

A comparison between EBU
and DQN

04

Conclusion

EBU and its challenges

Episodic Backward Update

- Sample a whole episode from replay memory
- Propagate the value for entire transitions in backward manner
- Diffusion factor is set to overcome overestimate error

```
7:      Sample a random episode  $E = \{S, A, R, S'\}$  from  $D$ , set  $T = \text{length}(E)$ 
8:      Generate a temporary target  $Q$ -table,  $\tilde{Q} = \hat{Q}(S', \cdot; \theta^-)$ 
9:      Initialize the target vector  $y = \text{zeros}(T)$ ,  $y_T \leftarrow R_T$ 
10:     for  $k = T - 1$  to  $1$  do
11:          $\tilde{Q}[A_{k+1}, k] \leftarrow \beta y_{k+1} + (1 - \beta) \tilde{Q}[A_{k+1}, k]$ 
12:          $y_k \leftarrow R_k + \gamma \max_a \tilde{Q}[a, k]$ 
13:     end for
14:     Perform a gradient descent step on  $(y - Q(S, A; \theta))^2$  with respect to  $\theta$ 
```

Implementation

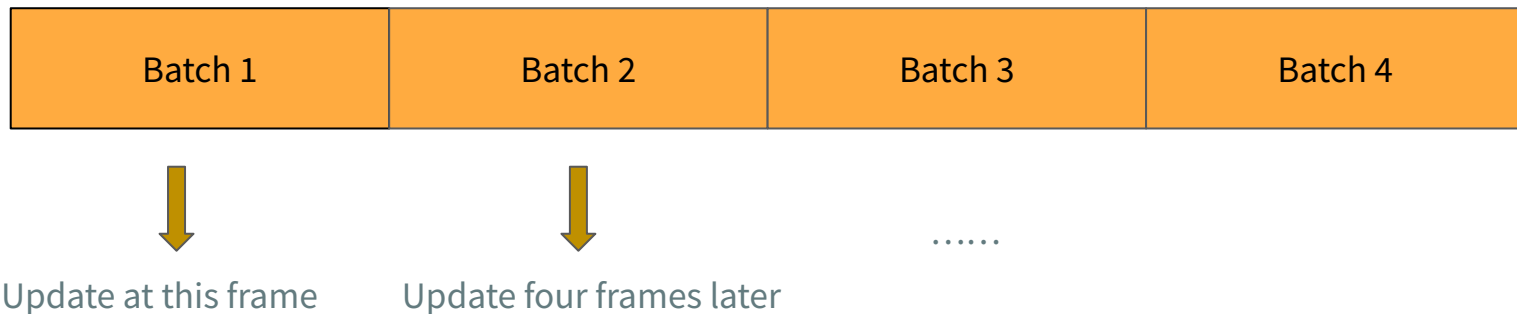
- Calculate $Q_{\tilde{}}$ with target network
- Generate y with $Q_{\tilde{}}$, actions, Beta, discount and rewards

```
if cur_frame % 4 == 0:
    if batchnum == batchcount:
        states, actions, rewards, next_states, dones = buffer.sample()
        transition_length = len(actions)
        batchnum = ceil(transition_length / batch_size)
        batchcount = 1
        target_net.eval()
        _Q = target_net(next_states) #(sample_length, num_action)
        target_net.train()
        _Q[-1,:] = torch.as_tensor([0] * num_of_act)
        y = np.zeros(transition_length)
        y[-1] = rewards[-1]
        for i in range(transition_length - 2, -1, -1):
            _Q[i][actions[i+1]] = B * y[i+1] + (1 - B) * _Q[i][actions[i+1]]
            y[i] = rewards[i] + discount * torch.max(_Q[i])
```

Implementation


- Sample a batch size of transitions and update with optimizer
- Sample another batch size from the same trajectory until encountering the termination flag

Whole episode from buffer: a trajectory



Results


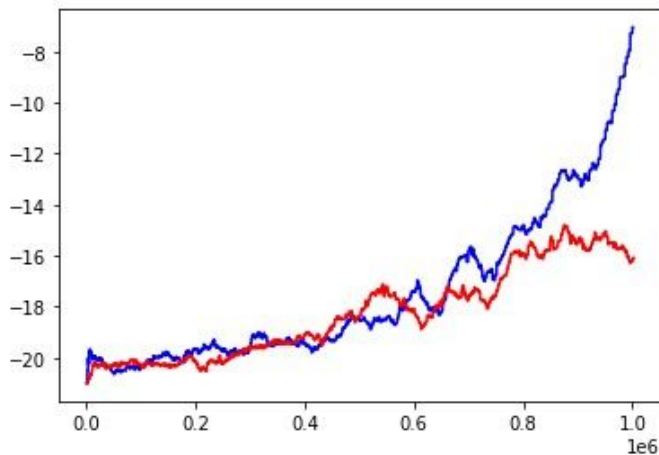
- Adam with low learning rate: $1e-5$
- $\beta = 0.5$
- Discuss:
 - 在低lr的時候EBU是能夠train得起來的
 - 在前 $1.5e6$ 的時候表現有稍微超過DDQN, 只是後面就輸掉了



Results



- RMSprop with higher learning rate: $2.5e-4$
- $\beta = 1$
- Trajectory saved with life loss, not termination
- Discuss
 - 前100萬個frame看起來競爭力都很夠
 - 100萬之後就會開始震盪上不去
 - Agent的表現很不穩定, 有時候last 25 eps的平均很低, 但是測試單局的分數又很高

Episode: 2880/5000, Epsilon: 0.100, Loss: 1.1884658, Return: -18.60
current episode reward: 21.0




Other Results

- Trying different parameters




Behavior


- DDQN



- EBU



Final experiment



Conclusion

$$\tilde{Q}[\mathbf{A}_{k+1}, k] \leftarrow \beta \mathbf{y}_{k+1} + (1 - \beta) \tilde{Q}[\mathbf{A}_{k+1}, k]$$

- Diffusion factor 設為一致的效果在Pong的環境中的表現較DQN差
- adaptive 版本的表現比其他model好(用不同的 β 去train很多個target network, 再選表現較好的network去update), 如果是固定的 β 去做training, model能train到的程度感覺很吃運氣, 很容易因為sample到的trajectory不夠好讓loss重新掉回好幾個eps前的表現
- 會說sample到的trajectory影響對model影響很大是因為我們有把agent evaluation的video 定期拿出來看, 常常原本agent有學會接球, 但過了幾個eps後又變成不會接發球, 直接卡在最上方或最下方, 我們推測是因為可能取到太多次發球在同一邊的trajectory, 造成model在遊戲一開始的對策一直變動



Thanks for Listening!